

Price Adjustment and Retail Information

Hong Kong Main Board IPOs in 2026

Prepared on 4 October 2026. English research report for supervisor discussion.

1. Findings, sample and definitions

We study the existing 113 ordinary IPOs listed from 2 January to 30 September 2026. The main price outcome remains offer-to-first-day-close return. Opening and intraday returns locate price changes. The existing HSI adjustment checks market sensitivity.

Only 43 offers have a usable two-sided filing range. The old field labels 70 offers “Fixed price”, but 31 of them lack a disclosed lower endpoint. The 39 equal-endpoint offers and the 31 missing-lower offers are separate here. We do not invent a midpoint or assign zero revision to either group. Original workbooks and extraction records were not changed. This study corrects the research classification, not source fields.

The revision–close association is positive but uncertain. A day-stage association is visible in some designs, but it does not pass the main adjusted inference. Higher price return also does not imply higher application return. Pre-deadline margin coverage is too selected for a general prediction claim.

Table 1. Coverage and returns

Price evidence	N	Mean IR	Median IR	APR	Ceiling APR
Two-sided price range	43	68.29	37.53	2.75	2.28
Equal price endpoints	39	73.39	44.74	2.41	2.41
Lower bound not disclosed	31	6.71	2.41	0.75	0.75

All return columns are percent. IR uses the offer price and first-day close. APR uses requested units at the final offer price. Ceiling APR uses the published maximum price. These are allocation expectations, not observed account returns.

Table 2. Within the 43 range-priced offers

Final position	N	Mean IR	Mean open	Mean APR	Median APR
At low	13	51.65	53.73	6.40	0.05
Inside range	13	62.19	57.47	-0.03	0.15
At high	17	85.67	77.09	2.07	1.00

No offer in this range sample is outside its recorded range. Group means are descriptive. Mean intraday return is not the difference between mean close and opening returns.

Let P_L, P_H be the filed endpoints, $P_M = (P_L + P_H)/2$, P_0 the final offer price, P_O the opening price and P_C the first-day close. For range offers,

$$v = P_0/P_M - 1, \quad z = (P_0 - P_L)/(P_H - P_L), \quad w = (P_H - P_L)/P_M.$$

Let $a = E[\text{allotted units}]/\text{requested units}$. The existing gross application return is $a(P_C/P_0 - 1)$. The maximum-price principal measure is $a(P_C - P_0)/P_H$. Both exclude brokerage, levies, fees, borrowing and tax. The latter uses a price known at launch as its principal base; it is not a measured account payoff. For all 113 offers, these two application means are 2.08% and 1.91%, respectively.

2. Price adjustment across the trading stages

The test asks whether a higher final price relative to the filed midpoint accompanies a higher later price return. Hanley (1993) motivates this partial-adjustment association [1]. It does not by itself identify bookbuilding information, intentional underpricing, or market inefficiency. The 43 range offers include 16 Q1, 14 Q2 and 13 Q3 IPOs, with nine A+H issuers. No new sector fixed effects are fitted to this small sample.

$$y_i = \alpha + \beta v_i + \delta_2 Q2_i + \delta_3 Q3_i + \varepsilon_i.$$

Each outcome uses exactly the same firms and controls. Log outcomes are used to reduce the influence of very large gains; raw close return remains reported. Allocation and application return enter on their original fractional scale.

Table 3. Common-sample regressions with quarter controls

Outcome	Slope	HC3 SE	HC3 p	WCR p	Holm WCR p
Log offer-to-close	15.978	13.723	0.251	0.207	1.000
Log offer-to-open	8.745	10.984	0.431	0.375	1.000
Log open-to-close	7.233	3.583	0.050	0.051	0.355
Allocation fraction	-0.250	2.970	0.933	0.781	1.000
Application return	-3.244	2.830	0.259	0.324	1.000
Log HSI-adjusted close	15.536	13.799	0.267	0.242	1.000
Raw offer-to-close	21.388	25.018	0.398	0.324	1.000

N=43, nine listing-month clusters, intercept and Q2/Q3 indicators in every column. Slope and SE are scaled for a +10 percentage point revision: log points for log outcomes, percentage points for fractions and raw returns. HC3 uses a t reference. WCR is restricted wild cluster-t with all 512 Rademacher sign vectors. Holm adjustment covers all seven outcomes in this table. With only nine month clusters, inference remains fragile. No significance stars are used.

A +10 pp revision has an estimated 17.33% multiplicative change in the close/offer ratio, not a 17.33 pp increase in raw return. Its uncertainty is large (Table 3). None of the seven main outcome tests survives Holm adjustment. The unadjusted day-stage p-values are near 0.05; they do not establish a robust mechanism.

Table 4. All four declared control designs

Controls	Log close slope	Log day slope	APR slope
Revision only	13.958	7.270	-3.243
Quarter controls	15.978	7.233	-3.244
Launch size and A+H	15.024	7.685	-3.418
Subscription HSI and launch size	16.790	7.856	-3.252

Same 43 IPOs in every design; slopes have the Table 3 scaling. These are alternative specifications, not a nested sequence. All 28 focused estimates and all coefficients are in the CSV output. Subscription HSI runs from prospectus-day close to subscription-day close; it is an explanatory market-window check, not verified actual-pricing-date information.

The price stages obey exact accounting:

$$\ln(P_C/P_0) = \ln(P_O/P_0) + \ln(P_C/P_O), \quad \ln(P_C/P_M) = \ln(P_O/P_M) + \ln(P_C/P_O).$$

The first identity makes the close slope the sum of the opening and day slopes. The second is a shared-price warning. Using $\ln(P_O/P_M)$ as regressor with the same quarter controls gives a midpoint-close slope that is exactly one plus the offer-close slope. The two estimates are 2.494 and 1.494; their difference is accounting, not an independent test.

3. Sensitivity and pre-deadline demand information

Table 5. Quarter and large-winner sensitivity

Retained sample	N	Log close slope	Log day slope	APR slope
All range offers	43	15.978	7.233	-3.244
Remove Q1	27	12.681	7.223	-4.443
Remove Q2	29	13.662	6.789	-3.211
Remove Q3	30	23.773	7.877	-1.651
Remove three largest IRs	40	15.933	7.676	-2.589
Remove three largest APRs	40	27.925	9.996	0.110

Each row uses the available quarter indicators. Exclusions retain identical firms across outcomes. Removing return leaders and removing application-return leaders answer different questions. Full leave-one-IPO-out results are also saved.

Single-issuer exclusions retain positive log-close slopes, from 10.231 to 23.398 log points per +10 pp revision. Application slopes remain negative in single exclusions, but turn positive when the three largest application-return winners are removed. Thus, the mean application pattern is sensitive to a small group of large payoffs. The lower-range group is not established as a superior investment rule.

Final demand, offer price and allocation are later outcomes. The margin exercise uses only snapshots published on a date strictly before the subscription closing date. Of 93 source observations, 78 meet that condition, 13 are after the deadline, and two have unresolved closing-day timing. We use the latest eligible observation for each covered issuer. This leaves 25 IPOs across six listing months; 14 are in July. The survey amounts cover named brokers and unspecified others, not the full market.

Table 6. Pre-deadline margin information: selected 25-IPO sample

Outcome	Slope	HC3 SE	HC3 p	WCR p	Holm WCR p
Log offer-to-close	14.345	12.105	0.249	0.438	0.875
Allocation fraction	-5.351	4.605	0.258	0.125	0.375
Ceiling-principal return	1.814	1.741	0.309	0.438	0.875

The regressor is $\ln(\text{latest safely published margin multiple})$; intercept and launch size are included. Slope units are 100 times the coefficient: log points or percentage points per one log unit of margin. $N=25$; six month clusters; WCR enumerates 64 signs. Holm covers the three controlled outcomes. Six clusters give coarse inference. The broker-survey scope is incomplete and coverage is selected.

More reported margin demand accompanies higher price returns in the point estimates and lower allocation fractions. Its application-return estimate is also positive, but uncertain. None of the controlled tests passes Holm adjustment. These associations cannot show that borrowing causes demand or that the signal predicts a general IPO cohort.

Table 7. Data use and information timing

Data	Meaning and permitted use
Filed range and global quantity	Launch terms; measure price uncertainty and planned deal scale.
Final offer price	Pricing outcome; construct revision and return, not a presumed deadline feature.
Final requests and allocation	Allotment outcomes; measure rationing and expected payoff, not advance signals.
Published margin snapshot	Partial broker demand before deadline; requires timestamp and survey-scope checks.
Opening and closing prices	Later outcomes; locate price changes and label training returns only after listing.
Recorded pricing date	Often expected or proxy date; cannot certify the actual information cutoff.

Range and planned-quantity inputs match the stored extraction values in all 113 IPOs: 308 non-missing field comparisons and 31 missing lower endpoints. The audit saves quotes, source pages, disclosure URLs and extraction hashes. A matching value does not replace independent semantic review of the disclosure. For example, OmniVision’s prospectus expressly permits disclosure of the maximum price only [2].

Only 36 of the 43 range offers have a recorded pricing-date proxy; seven lack that anchor. These are not verified actual pricing dates. We therefore use the complete subscription-to-listing HSI window for the return check. This prevents date availability from changing the main firm sample, but it does not measure the true bookbuilding window.

4. Temporal assessment and research conclusion

We also test a small model using two launch variables: log planned proceeds and A+H status. Planned proceeds use the filed global quantity times the midpoint for range offers, or the published ceiling when no two-sided range exists. Final proceeds, final demand and final allocation do not enter the forecast.

For each target, training IPOs must have completed their first trading day before the target prospectus date. Each prediction requires at least 40 such IPOs. This gives 70 evaluated targets. Outcomes of later or same-day listings cannot enter training. An expanding historical mean is the benchmark. No allocation strategy is simulated.

Table 8. Retrospective temporal prediction check

Model	N	MAE (pp)	RMSE (pp)	MSE gain (%)
Expanding historical mean	70	3.398	5.963	0.00
Launch size and A+H	70	3.345	5.915	1.62

Target: gross expected return on maximum-price requested principal. Each target uses at least 40 prior IPOs whose first-day close precedes its prospectus date. The regression uses only filed offer size and A+H status. MSE gain is relative to the expanding-mean forecast. The current sample was already explored; this is not a new untouched holdout or a trading backtest.

The size-and-A+H model reduces mean squared prediction error by only 1.62%. This small retrospective gain does not establish useful advance selection. Mean actual ceiling-principal return in the evaluated subset is 2.03%. Forecast errors are large relative to that mean. The exercise uses already explored data and offers no untouched test of future performance.

Conclusion. The current data do not establish a strong partial-adjustment mechanism in Hong Kong’s 2026 range offers. Positive share-return slopes coexist with weak allocation evidence and winner-sensitive application returns. The available pre-deadline broker data do not yet support a reliable selection claim.

Next evidence work. First verify the actual price-determination dates and the maximum-only status against disclosures. Next expand timestamped margin coverage across months and document the same broker scope. Then freeze a parsimonious model before collecting a later, untouched IPO cohort. Test price return, allocation and application-principal return separately. Use observed costs only in a later net-return exercise. Grey-market trading after subscription closes cannot be an application signal.

Supervisor discussion draft

The earlier study found a large gap between first-day share returns and retail application returns. This follow-up tests the pricing stage and the information available before applications close. Of 113 IPOs, 43 disclose usable price ranges; 31 other offers were incorrectly grouped with fixed-price offers when their lower bound was missing. Within the common range sample, price revision has a positive but uncertain association with first-day close return. Application-return results depend on a few large winners. Timestamped margin data cover only 25 selected IPOs. A retrospective launch-information model offers little forecast improvement. We therefore retain the close return as the main outcome and prioritise timing and coverage checks before making an information-based or predictive claim.

Sources and reproduction

[1] Hanley, K. W. (1993). The underpricing of initial public offerings and the partial adjustment phenomenon. *Journal of Financial Economics*, 34, 231–250. Publisher DOI.

[2] OmniVision Integrated Circuits Group, prospectus dated 31 December 2025, waiver discussion; extraction PDF page 109 (printed page 101). HKEX disclosure.

[3] Project source ledger: prospectus PDFs, frozen retail tiers, raw first-day prices, HSI daily cache, and timestamped broker-survey observations. The report is an extension of these stored sources, not a new full-data certification.

Run `python run.py analysis -study pricing_adjustment` from the repository root.

Study output: `analysis/out/pricing_adjustment/`. The source audit, common samples, all models, exclusions, timing ledger and rolling training membership are saved. Input hashes and software versions are in `run_manifest.json`.